

# CSC 7700: Scientific Computing

## Module D: Simulations and Application Frameworks

### Lecture 1: Simulation Basics

Dr Peter Diener

Center for Computation and Technology  
Louisiana State University, Baton Rouge, LA

November 15, 2013



- 1 Goals
- 2 Simulations
- 3 Supercomputers



# Goals









- The module *Simulations and Application Frameworks* will teach:
  - what a simulation is,
  - how a typical simulation code looks like,
  - how it is used in practice,
  - and what some of the major concerns in such a code are.
- We will use Cactus as an example of an Application Framework.



# Simulations



# What is a simulation?




# What is a simulation?

- Flame propagation in combustion engine: *understand* behavior that is too fast or too small.
- Hurricane modeling: *predict* behavior.
- Car crash testing: *engineer* better devices.
- Video games: *create* a fantasy world.
- Black hole collisions: *test* the basic laws of nature.



# What is a simulation?

Laws of nature  
(Physics, Chemistry, Biology...)



Simulation





# What is a simulation?

Laws of nature  
(Physics, Chemistry, Biology...)



Mathematical  
model



Supercomputer



Simulation



# Ingredients of a simulation

- From the laws of nature derive a mathematical model relating the variables describing the system you want to simulate. This often involves approximations.
- Decide on a discretisation scheme to use for your variables and your mathematical equations (more later).
- Write a computer code to solve the discretised equations.
- Setup initial data and run the simulation.
- Analyze and visualize the results.



## Example: The wave equation

Consider a fluid in 1D described at time  $t$  and position  $x$  by the density  $\rho$ , pressure  $p$  and velocity  $u$ .

From Newton's first law ( $ma = F$ ) we have

$$\rho \frac{\partial u}{\partial t} = -\frac{\partial P}{\partial x}.$$

On the other hand if the velocity  $u$  varies with position  $x$  the pressure will change

$$\frac{\partial P}{\partial x} = -K \frac{\partial u}{\partial x},$$

where  $K$  is the incompressibility of the fluid.

Assuming that  $K$  and  $\rho$  does not vary with position we can combine the equations into one equation for the pressure  $P$

$$\frac{\partial^2 P}{\partial t^2} = \frac{K}{\rho} \frac{\partial^2 P}{\partial x^2}.$$



# The Initial Value Problem

The wave equation is a very simple example of a very large class of computational problems encountered in physics, chemistry and biology: The Initial Value Problem (IVP).

- Typically expressed in terms of a set of PDEs (partial differential equations) that tells us how a system is changing given a known state of the system.
- Starting from a set of *Initial Conditions* we can then simulate the behavior of the system by evolving from one state to another in finite timesteps.



- It would require an infinite amount of information to describe the complete state of continuum system (air, water, car body, etc. ).
- Instead we select a finite set of discrete points where we assign values to the field variables in order to reduce the number of degrees of freedom.
- There are many different ways to do this
  - Finite differences (sample solution on a regular grid)
  - Finite elements (small rigid triangles or tetrahedra)
  - Finite volumes (small squares or cubes)
  - Expansion of fields in a finite number of basis functions.
  - Particles




- Using a finite (instead of infinite) number of degrees of freedom introduces an approximation error.
- The accuracy of the approximation depends on the discretisation scheme and the resolution (how closely is the continuum sampled).
- For any discretisation scheme we can establish the order of accuracy. E.g. fourth order:  $\epsilon(h) = O(h^4)$ .
- Validation of simulation results requires careful convergence studies.
- The order of accuracy only gives a reliable measure of the error when the resolution is high enough (the convergent regime).



# Discretisation example: The wave equation

Approximating the pressure  $P(x, t)$  with a grid function  $P_i^{(n)}$ .




$$\begin{aligned}\frac{\partial^2 P}{\partial t^2} &= v^2 \frac{\partial^2 P}{\partial x^2} \\ \Downarrow \\ \frac{P_i^{(n+1)} - 2P_i^{(n)} + P_i^{(n-1)}}{\Delta t^2} &= v^2 \frac{P_{i+1}^{(n)} - 2P_i^{(n)} + P_{i-1}^{(n)}}{\Delta x^2}\end{aligned}$$

The error from this time and space discretisation is  $O(\Delta x^2)$ .



# Caveat

- Some systems are not described by PDEs but other types of mathematical relations.
- Sometimes it is the end state of a system that is interesting (like an equilibrium configuration) and not how it got there.
- If the initial data is an approximation or a guess, the result of a simulation may not be reliable.





# Simulation Overview

- Derive a mathematical model for the system you want to simulate.
- Discretise the resulting PDEs (or whatever other form they are in).
- Decide on the shape and size of the computational domain (normally it's impossible to simulate the whole universe).
- Decide on the resolution.
- Decide on how to handle the boundaries of the computational domain.
- Set up initial conditions.
- Evolve using many small timesteps.
- Output data along the way.
- Analyse and visualize the data.
- Write the paper.



# Structure of a Simulation Code

- Variables (solution) stored in large “state vectors”
  - There can be many (e.g. billions) elements in the state vector
  - Best handled in efficient container structure like a Fortran array, C or C++ vector or a tree structure.
  - The code needs constructs to iterate over these elements or subsets of elements.
- Routines to set up the initial condition.
- Routines to perform many identical time steps.
- I/O methods to write solution to disk and optionally checkpoint the simulation.
- Should be able to run as batch job without user intervention.



- Analytical initial data can just be calculated when the simulation starts.
- Numerical initial data may be solved internally at the beginning or be read in from file.
- Can be set up from reading in a checkpoint file (continuation run).
- Initial data can be large:
  - 1 billion elements,
  - 25 variables/element,
  - 8 Byte/variable: total 200 GByte.



# Parallel Computing

- Cannot store the solution on a single node; parallel programming using MPI is necessary.
- At the moment only Fortran, C and C++ are viable languages for programming a supercomputer.
- Research is underway in other (maybe simpler) ways like Unified Parallel C, Co-Array Fortran or HPX (here at LSU)
- Simulations take long:
  - 1 billion elements,
  - 1000 Flop/element/step,
  - 1 million steps,
  - CPU speed 10 GFlop/s: total 28,000 CPU hours (3.2 CPU years), or 12 days when running on 100 CPUs.



- Since simulations take so long we cannot supervise them manually
  - Cannot be awake at all times.
  - A user error can destroy weeks of data.
  - Supercomputers are expensive; cannot waste valuable resources waiting for the next user input.
- Need to *plan* simulations carefully ahead of time, then let them run automatically.



- Need to *plan* simulations carefully ahead of time, then let them run automatically...
- ...so that each error is only discovered days or weeks later!
- Using a supercomputer thus requires much expertise, experience, patience and a high tolerance for frustration.
- *Using supercomputers are more difficult than it really should be.*



# Supercomputers



# Fast vs. Large


- Supercomputers are not so much fast as they are large.
- They are not interactive (like a notebook or workstation), they operate in batch mode.
- Their hardware is complex.





# Remote Access

- Supercomputers are located far away, need to use ssh/gsissh to access.
- Log in is to *front end (head node)* only, usually a large workstation.
- Cannot (or should not) use front end to run simulations.



- Supercomputers need large and fast file systems to store simulation data often many 100 TByte or PByte.
- For management and performance reasons the file systems are usually split into different parts with different properties.

**Home directory** GBytes per user, many small files, backed up.

**Data directory** TBytes per user, few large files, backed up, tape backend.

**Scratch directory** No quota, few large files, often automatically deleted.

- This is done differently on each supercomputer – read documentation!



# Compute Nodes and Interconnect

- Most supercomputers have a *cluster* architecture with many interconnected compute nodes.
- Each node may have multiple cores (4 to 64), similar to a large workstation typically with 1GByte to 4GByte of memory per core.
- Nodes may additionally have GPU's or Intel Xeon Phi's.
- Nodes are connected via a low-latency *communication network* (e.g. Infiniband or other proprietary technologies).
- Overall system has from a few up to 10,000 nodes. The largest Supercomputer in the world currently has 3,120,000 cores in 16,000 nodes (Intel Xeon IvyBridge processors and Xeon Phi).
- My personal scale:
  - small < 2k cores
  - medium < 20k cores
  - large > 20k cores



# Batch System

- Cannot (or should not) use compute nodes directly.
- Need to submit *job* to *batch system*, requesting  $N$  nodes for  $T$  amount of time.
  - ...wait (a few days?)...
  - ...then the job start to run (most likely while you're asleep)...
  - ...and then you discover your errors...
- There is always a run time limit (e.g. 8, 12, 24, 48, 72 hours)
  - ...which is inconvenient when you need to run for 2 weeks:  
checkpoint/restart is necessary.
- The batch system ensures that jobs get run in an order that keeps the supercomputer as busy as possible at all times.



# Allocations

- Need to ensure fair use of supercomputers and prevent individual users from monopolising it.
- Need to ensure that supercomputer time are used for important research and by codes that can run efficiently.
- Allocation proposals for the national supercomputers are therefore peer reviewed.
- The allocation review panel then decides how much time each research project will be allowed to use in the coming allocation period.
- 1 CPU hour costs about 5 cents (10 cents on Amazon ECC).
- With this metric, Super Mike II produces about \$352 worth of CPU time every hour.



- Installed/available software is system dependent, not just standard Unix systems.
- Therefore cannot just install binaries, need to build software for each supercomputer.
- Installed software typically consists of compilers, MPI libraries, Scientific libraries (like GSL, BLAS/LAPACK, PETSC, etc.), perl, python, etc.
- Most interactions with a supercomputer takes place on the command line. Some GUIs, Portals exist, but typically limited in functionality.



# Supercomputer Steps

- Obtain an *Allocation*.
- Log in to the *Front End*.
- Compile your code on the *Front End*.
- Submit your jobs on the *Front End* to the *Batch Queue*.
- Simulations then execute on the *Compute Nodes* connected via a *Communication Network*.
- Data stored in *File Systems* can then be analyzed on the supercomputer itself or transferred to other places for analysis.
- Data can be stored permanently in *Tape Archives*.

