

Module E: Distributed Scientific Computing

Introduction to M-W Pattern,
MapReduce and Cloud Computing

Dr Shantenu Jha

<http://radical.rutgers.edu>

Overview of Module E

Distributed Scientific Computing

- Introduction to M-W and Cloud Computing
 - Master-Worker Pattern
 - Examples of M-W Pattern:
 - M-W Example Using SAGA: Mandelbrot Set
 - Ensemble simulations, Replica-Exchange
 - Introduction to MapReduce
 - Wordcount using SAGA MapReduce
- Introduction to Cloud Computing
 - Why Cloud Computing?
 - Convergence of multiple trends: Data-centric, Data-Center...
 - Understanding Amazon EC2 – default ‘standard’

Master-Worker Pattern

- Pattern: A commonly occurring mode of computation
 - Multiple patterns exist
 - e.g., publish-subscribe, broker etc.,
 - But M-W arguably one of the most pervasive
- M-W: Used in parallel and distributed computing
 - Simply put: Master assigns task to a worker; worker does work; gets more from Master
 - Master coordinates task distribution
 - M-W not an application in of itself, but a programming model or “communication pattern” upon which applications can be built
- What types of tasks are suitable for M-W?
 - Many independent “units” of loosely coupled tasks
 - Concurrent execution is feasible/permissible

Master-Worker Pattern

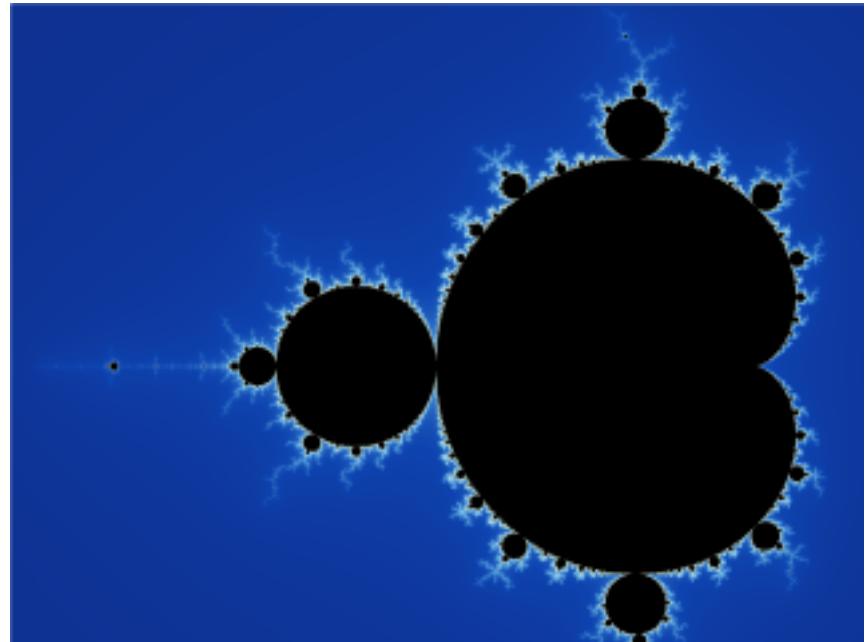
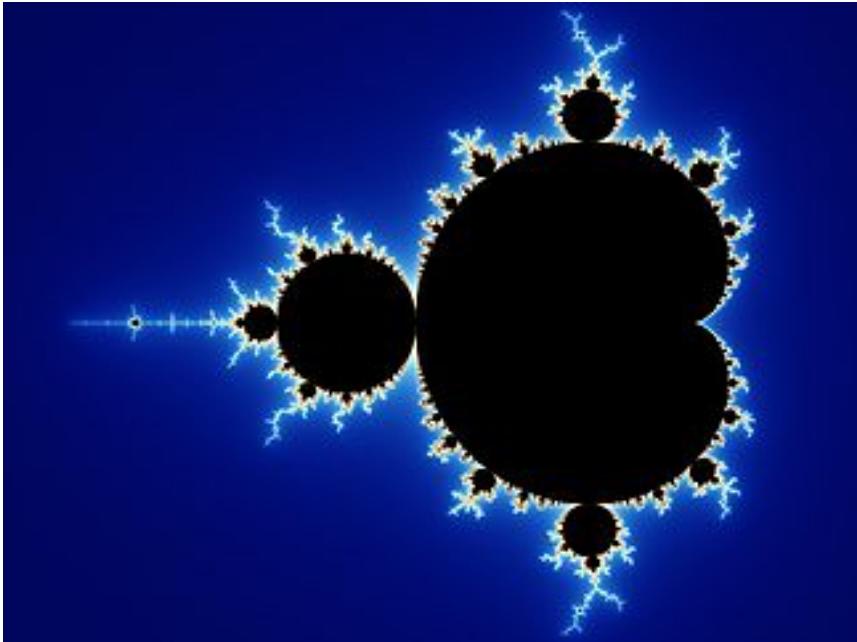
- What types of tasks are not suitable for M-W?
 - Decomposing into smaller independent units is not trivial
 - Lots of communication:
 - Either between Master and a Worker (s)
 - Master becomes the bottleneck!
 - Or between workers?
- Of Applications in E1, which are/can be M-W?
 - Nektar? Montage? SCOOP? Climateprediction.net?
 - Ensemble simulations and/or Replica-Exchange

Some Challenges in Distributed M-W Execution

- Task Decomposition and coordination:
 - How do we assign work units to workers?
 - What if we have more work units than workers?
- Execution and Fault-Tolerance:
 - How do we know all the workers have finished?
 - What if workers die?
- Coordination:
 - What if workers need to share partial results?
 - How do we aggregate partial results?
- Q: Based upon the above, what other constraints on suitability for M-W?

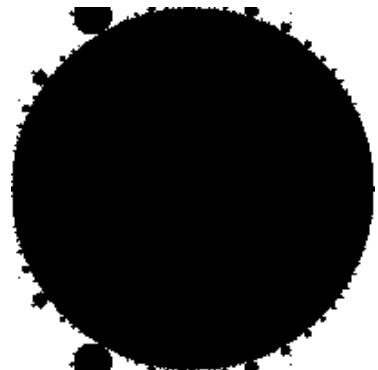
MANDELBROT SET

M-W to Compute Mandelbrot Set



How is M-W used to compute Mandelbrot Set?

- Task item: Complex plane broken-up;
compute parts of it
- Master puts task items into bucket.
Worker collects tasks;



SAGA-Based M-W: Mandelbrot

- You've seen Mandelbrot using SAGA-Python and BigJob
- Q: Discuss the similarities and the differences?
Are they both implemented as a M-W pattern?

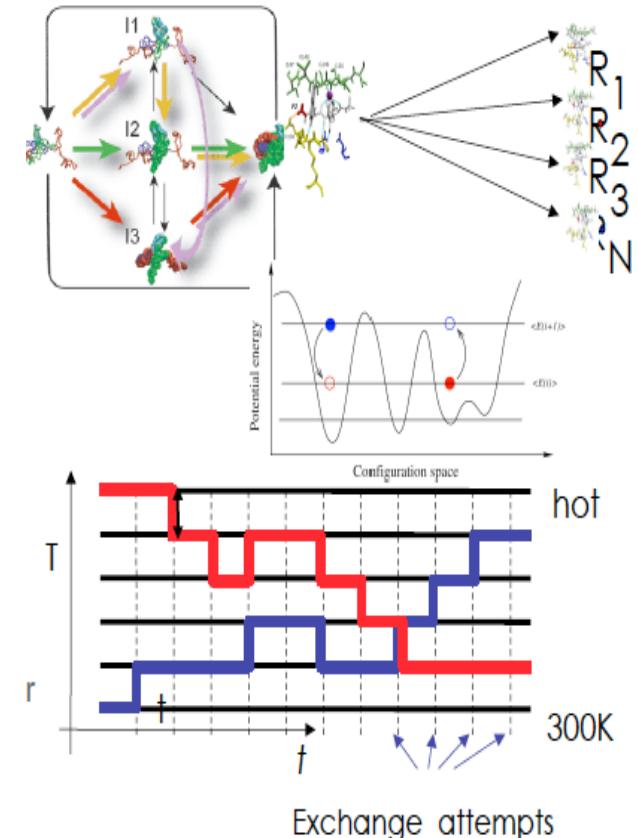
SAGA-Based M-W: Mandelbrot

1. Everything local: For 1 Master and same workload vary: $N_w = 2, 4$ and 8 Plot times to completion.
2. Distribute (equally?) the workers across a couple of XSEDE machines. Compare with (i) and (2)

ENSEMBLE-BASED REPLICA-EXCHANGE

Ensemble-based & Replica-Exchange Simulations

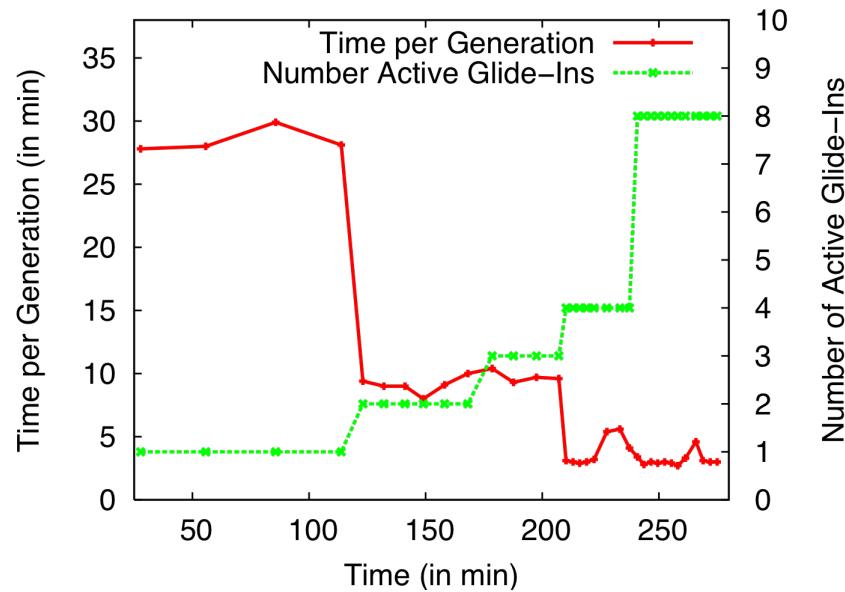
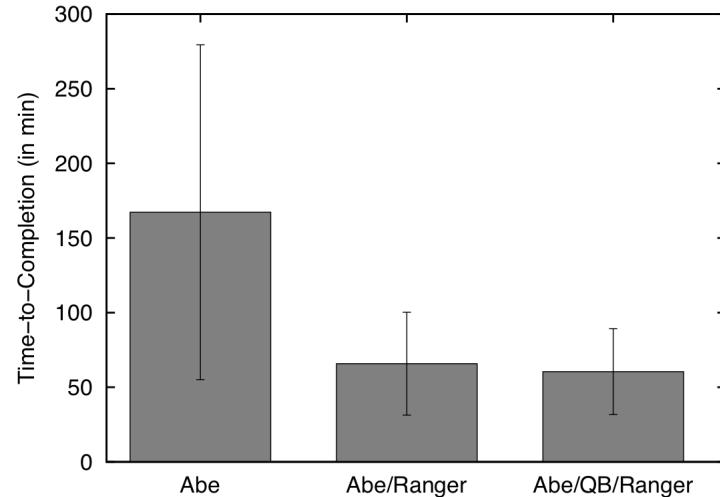
- Ensemble-based:
 - Many uncoupled simulations
 - But not necessarily uncoupled in analysis!
- Replica-Exchange (RE) methods:
 - Represent a class of algorithms that involve a large number of loosely coupled ensembles.
- RE simulations are used to understand a range of physical phenomena
 - Protein folding, unfolding etc
 - MC simulations
- Many successful implementations
 - Eg folding@home [replica based]



Distributed Adaptive Replica Exchange (DARE)

Multiple Pilot-Jobs on the “Distributed” TeraGrid

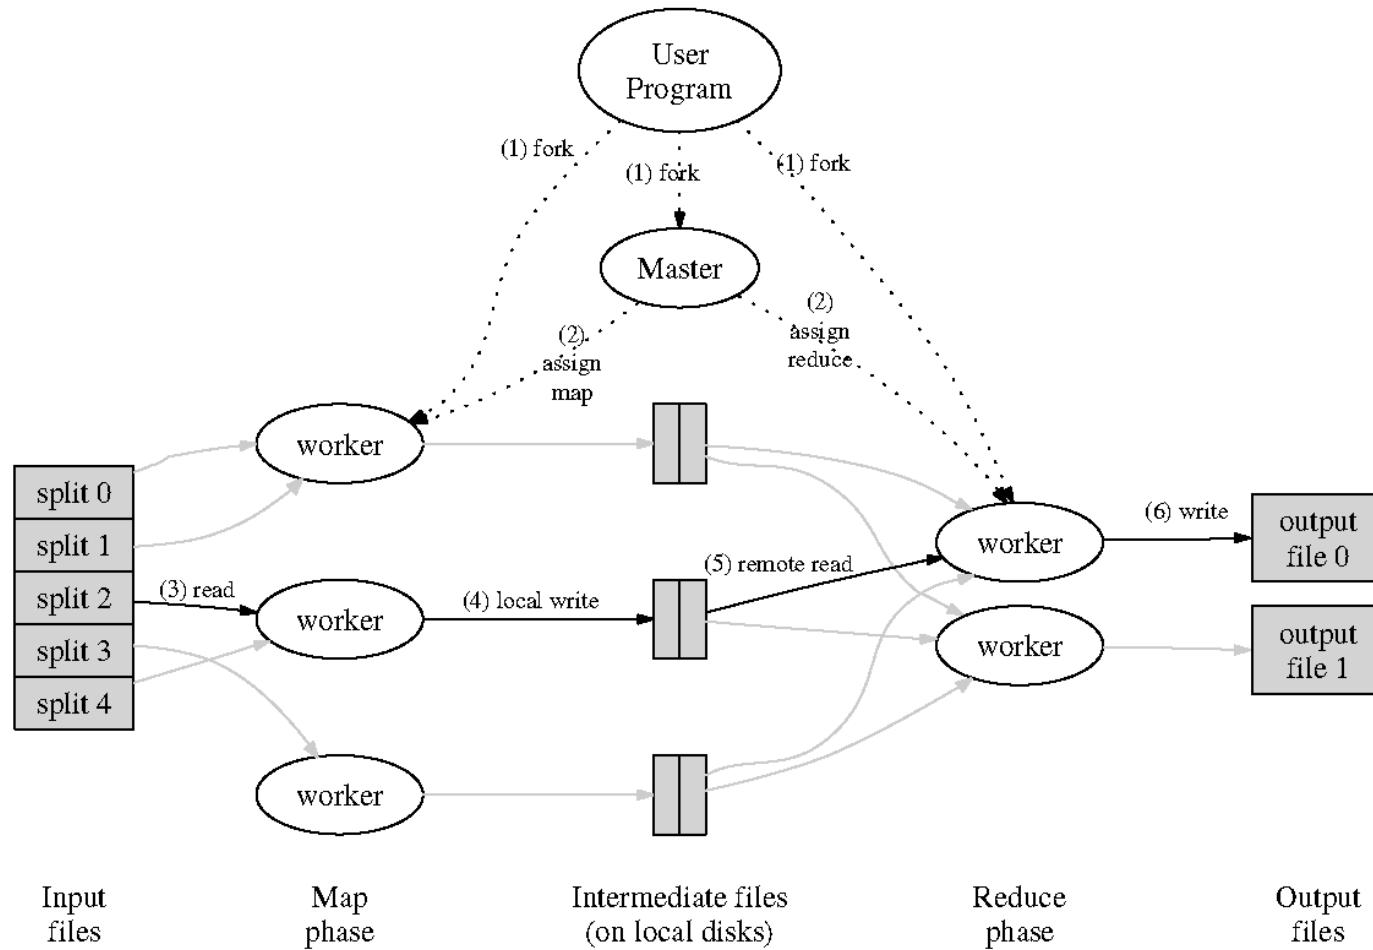
- Ability to dynamically add HPC resources. On TG:
 - Each Pilot-Job 64px
 - Each NAMD 16px
- Time-to-completion improves
 - No loss of efficiency



Understanding Replica-Exchange

- Why Distributed?
 - Many un-coupled units (ensembles/replica)
 - More resources, the merrier!!
- How Distributed?
 - Many implementations exist (eg folding@home)
 - SAGA-based “Pilot-Jobs” to use many distributed TG resources
- Limitations and Success?
 - Getting SAGA working on all machines!
 - Finding the best set of resources
 - Coordinating work across all the resources

MAP-REDUCE

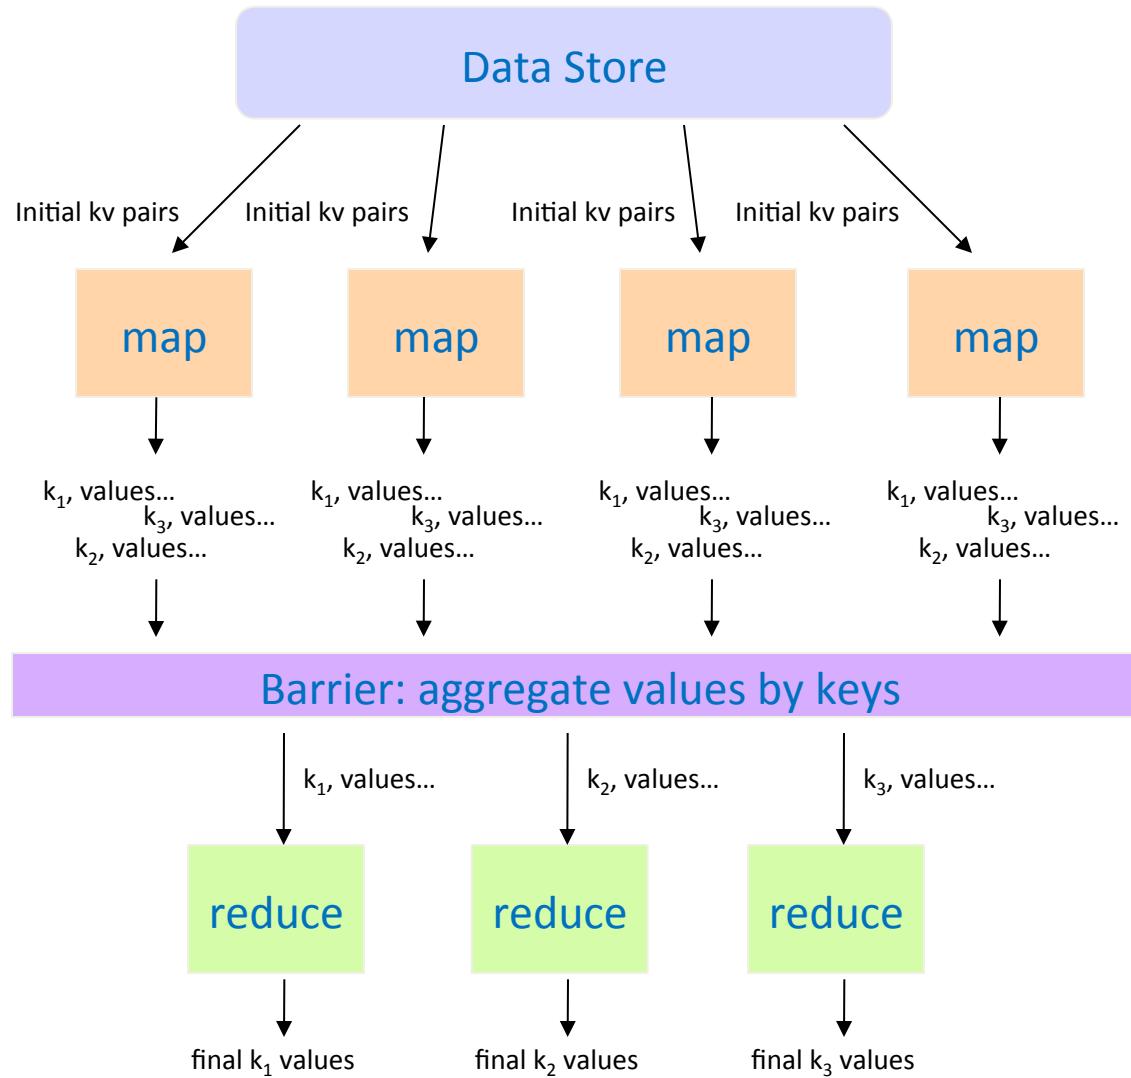


“Hello World”: Word Count

```
Map(String input_key, String input_value):  
    // input_key: document name  
    // input_value: document contents  
    for each word w in input_values:  
        EmitIntermediate(w, "1");
```

```
Reduce(String key, Iterator intermediate_values):  
    // key: a word, same for input and output  
    // intermediate_values: a list of counts  
    int result = 0;  
    for each v in intermediate_values:  
        result += ParseInt(v);  
    Emit(AsString(result));
```

Word Count via MapReduce



Some Challenges in Distributed M-W Execution (Redux)

- Task Decomposition and coordination:
 - How do we assign work units to workers?
 - What if we have more work units than workers?
- Execution and Fault-Tolerance:
 - How do we know all the workers have finished?
 - What if workers die?
- What if workers need to share partial results?
- How do we aggregate partial results?

MapReduce versus Google MapReduce (Runtime)

- MapReduce the pattern versus MapReduce the runtime
- Handles scheduling
 - Assigns workers to map and reduce tasks
- Handles “data distribution”
 - Moves the process to the data
- Handles synchronization
 - Gathers, sorts, and shuffles intermediate data
- Handles faults
 - Detects worker failures and restarts
- Everything happens on top of a distributed FS (later)

WORDCOUNT USING SAGA MAPREDUCE

SAGA MAPREDUCE

- Not tied to any specific infrastructure
 - Interoperable across different back-ends
 - No optimization, thus performance barrier
- Can control chunk-size, task size granularity, decomposition and placement/distribution
- Master-Worker pattern
 - Uses Advert Service to coordinate and distribute
- Contrast with Google MapReduce or Hadoop
 - Google/Yahoo extensively use the File-System
 - SAGA's flexibility comes at a performance!

SAGA PMR

- SAGA-based (Pilot) MapReduce:
 - <https://github.com/saga-project/PilotMapReduce>

SAGA PMR: WORDCOUNT

- Word Count:
 - <https://github.com/saga-project/PilotMapReduce/tree/master/applications/wordcount>
- *Generate your own input file for the wordcount example*
 1. Everything local: For 1 Master and same workload vary: $N_w = 2, 4$ and 8 Plot times to completion.
 2. Distribute (equally?) the workers across a couple of XSEDE machines. Compare with (i) and (2)
 3. Describe the role of the Pilot-Job?



Cloud Computing

What is cloud computing?



- *I don't understand what we would do differently in the light of Cloud Computing other than change the wordings of some of our ads*

Larry Ellision, Oracle's CEO

- *I have not heard two people say the same thing about it [cloud]. There are multiple definitions out there of “the cloud”*

Andy Isherwood, HP's Vice President of European Software Sales

- *It's stupidity. It's worse than stupidity: it's a marketing hype campaign.*

Richard Stallman, Free Software Foundation founder

What is a Cloud?

From NIST

- ***Resource pooling.*** Computing resources are pooled to serve multiple consumers.
- ***Broad network access.*** Capabilities are available over the network.
- ***Measured Service.*** Resource usage is monitored and reported for transparency.
- ***Rapid elasticity.*** Capabilities can be rapidly scaled out and in (pay-as-you-go)
- ***On-demand self-service.*** Consumers can provision capabilities automatically.

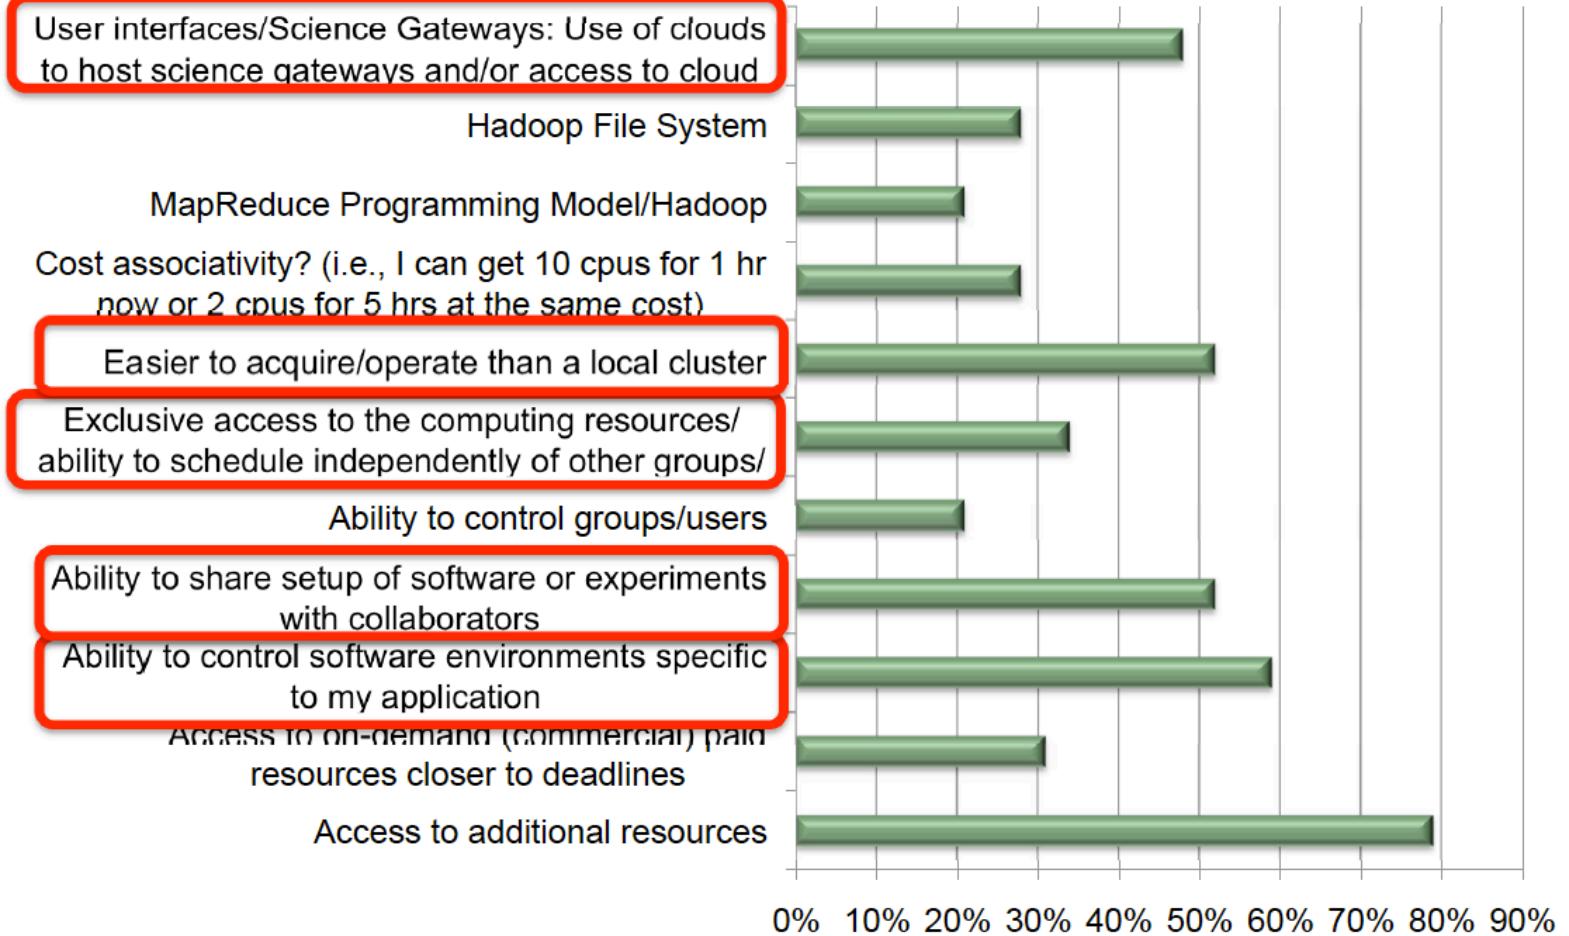
Why is This not Good Ol' Supercomputing

- A Supercomputer is designed to scale a single application for a single user.
 - Optimized for peak performance of hardware.
 - Batch operation is not “on-demand”.
 - Reliability is secondary
 - If MPI fails, app crashes. Build checkpointing into app.
 - Most data center apps run continuously (as services)
- Yet, “in many ways, supercomputers and data centers are like twins separated at birth.”*



Cloud Computing Interest

(adapted from Kathy Yelick)



Cloud Models

Infrastructure as a Service

- Provide a way to host virtual machines on demand
 - Amazon ec2 and S3 – you configure your VM, load and go

Platform as a Service

- You write an App to cloud APIs and release it. The platform manages and scales it for you.
- Google App engine:
 - Write a python program to access Big Table. Upload it and run it in a python cloud.
 - Hadoop and Dryad are application frameworks for data parallel analysis

Software as a Service

- Delivery of software to the desktop from the cloud
 - Stand-alone applications (Word, Excel, etc)
 - Cloud hosted capability
 - doc lives in the cloud
 - Collaborative document creation
- For more details on *aaS see paper by Lamia Youssef and Rich Wolski (GCE'09 @ SC09)

Cloud Computing: Enabling Technologies

(adapted from Kathy Yelick)

- **Centralization to lower costs**
 - Cheaper power due to bulk rates
 - Cheaper hardware purchase
 - Personnel savings from scale
- **Virtualization**
 - Allows sharing of resources
 - Allows tailoring software
- **Simple programming/usage models**
 - Preinstalled software services
 - Map-reduce

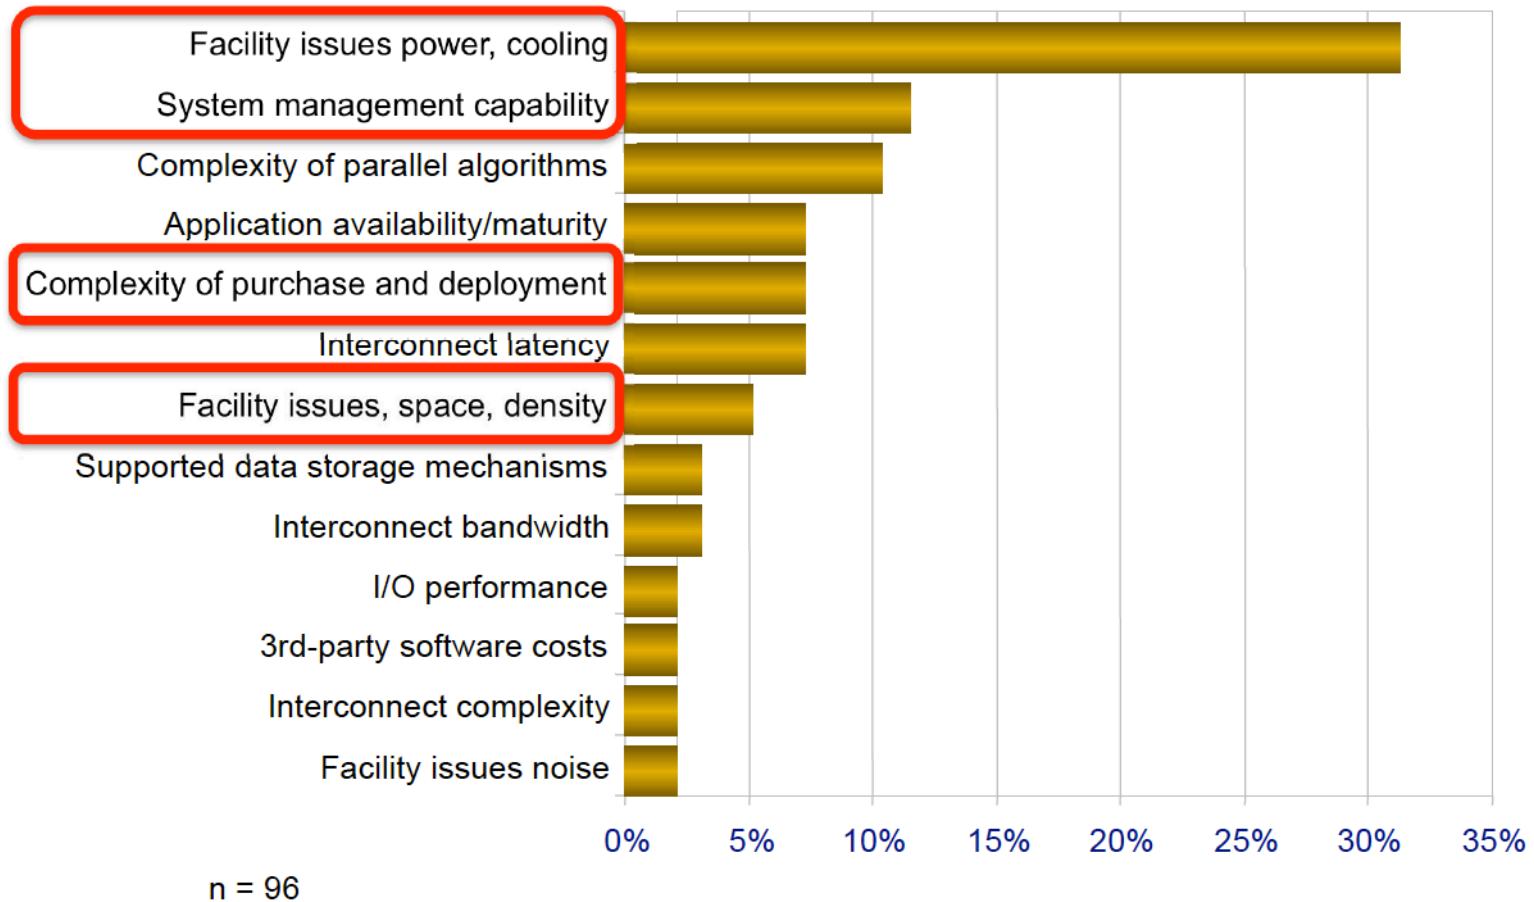
.. Its all about the Buisiness Model

(adapted from Kathy Yelick)

- **Cloud computing is a business model**
- It can be used on HPC systems as well as traditional clouds (ethernet clusters)
- Can get on-demand elasticity through:
 - Idle hardware (at ownership cost)
 - Sharing cores/nodes (at performance cost)
- How high a premium will you pay for it?
- How predictable is your workload?
 - Are data-intensive loads more predictable?

Top challenges to running own cluster

(adapted from Kathy Yellick)



The Data Center Landscape

Range in size from “edge” facilities to megascale.

Economies of scale

Approximate costs for a small size center (1K servers) and a larger, 50K server center.



Technology	Cost in small-sized Data Center	Cost in Large Data Center	Ratio
Network	\$95 per Mbps/month	\$13 per Mbps/month	7.1
Storage	\$2.20 per GB/month	\$0.40 per GB/month	5.7
Administration	~140 servers/Administrator	>1000 Servers/Administrator	7.1

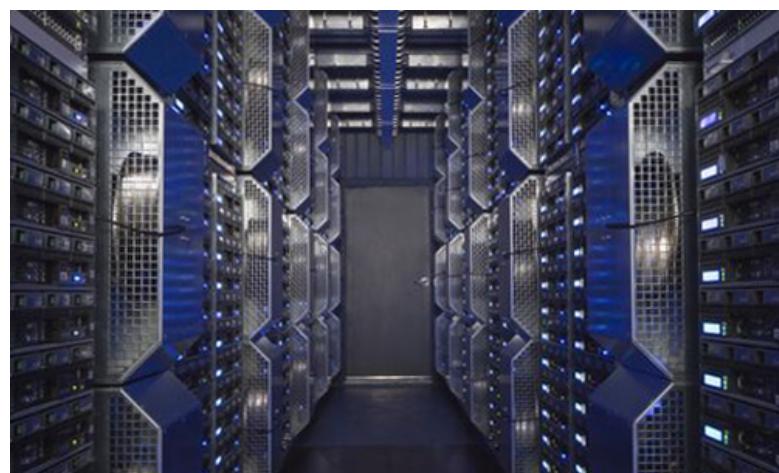
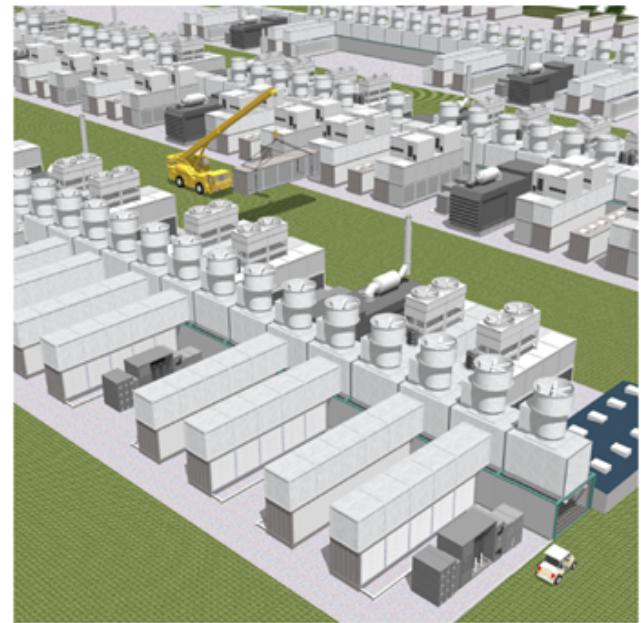


Each data center is
11.5 times
the size of a football field

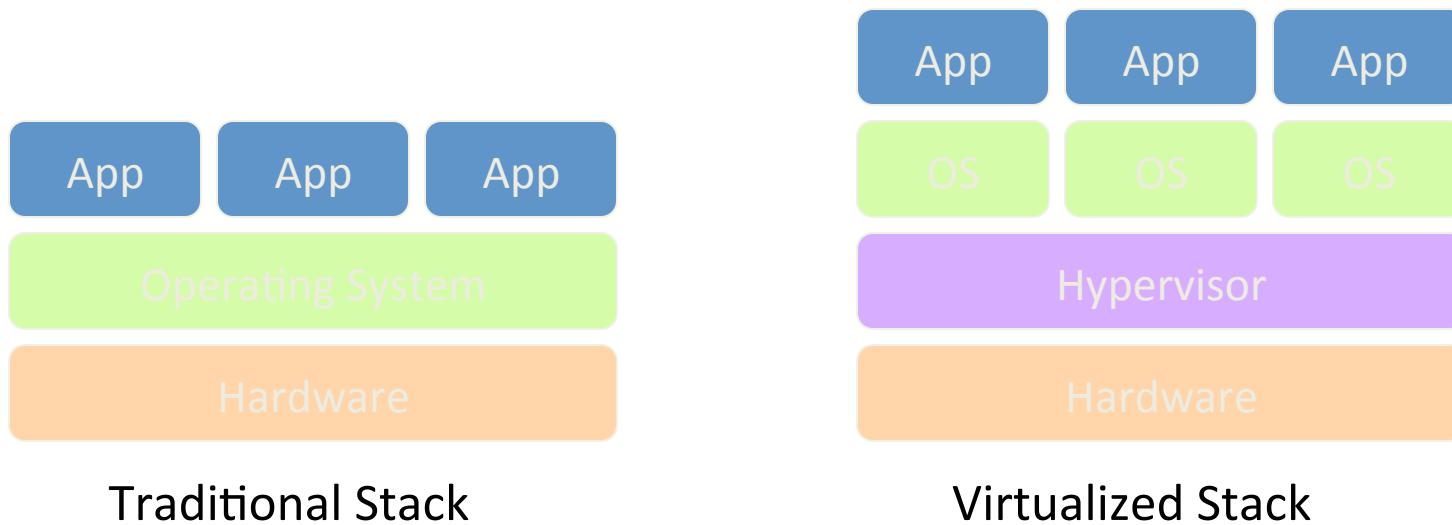
Advances in DC deployment

Conquering complexity.

- Building racks of servers and complex cooling systems all separately is not efficient.
- Package and deploy into bigger units:



Key Technology: Virtualization



Amazon AWS

<http://aws.amazon.com>

- Story goes: Build capacity for X-mas. What do with spare capacity year around?
- “Utility Computing”
 - Around long before Amazon EC2
 - \$0.10 per CPU-hour, plus bandwidth cost
- *aaS Model:
 - * = Infrastructure, Software, almost anything
- AWS: A set of APIs which give users access to Amazon technology and content
 - IaaS, but also “people as a service” – Mechanical Turk

Amazon Simple Storage Service (S3)

- Data Storage in Amazon Data Center
- Web Service interface
- No set-up fee, No monthly minimum
- Storage: \$0.15 per GB/Month
- Data Transfer: \$0.20/GB to transfer data
- Private and public storage
- Each object up to 5GB in size

Amazon Elastic Compute Cloud

- A Web service that provides resizable compute capacity in the cloud. Designed to make Web-scale computing easier
- A simple Web service interface that provides complete control of your computing resources
- Quickly scales capacity, both up and down, as your computing requirements change
- Changes the economics of computing:
 - Pay only for capacity that used; no cost of ownership
 - $a + bc$ becomes just bc

Amazon Elastic Compute Cloud

- No start-up, monthly, or fixed costs
 - \$0.10 per CPU hour
 - \$0.20 per GB transferred across Net
- No cost to transfer data between Amazon S3 and Amazon EC2
- More when we do Cloud Computing next week

Amazon Web Services

Default “community” standard

- Compute
 - Elastic Compute Service (EC2)
 - Elastic MapReduce
 - Auto Scaling
- Storage
 - Simple Storage Service (S3)
 - Elastic Block Store (EBS)
 - AWS Import/Export
- Messaging
 - Simple Queue Service (SQS)
 - Simple Notification Service (SNS)
- Database
 - SimpleDB
 - Relational Database Service (RDS)
- Content Delivery
 - CloudFront
- Networking
 - Elastic Load Balancing
 - Virtual Private Cloud
- Monitoring
 - CloudWatch

<http://aws.amazon.com/>

Elastic Compute Cloud (EC2) Service

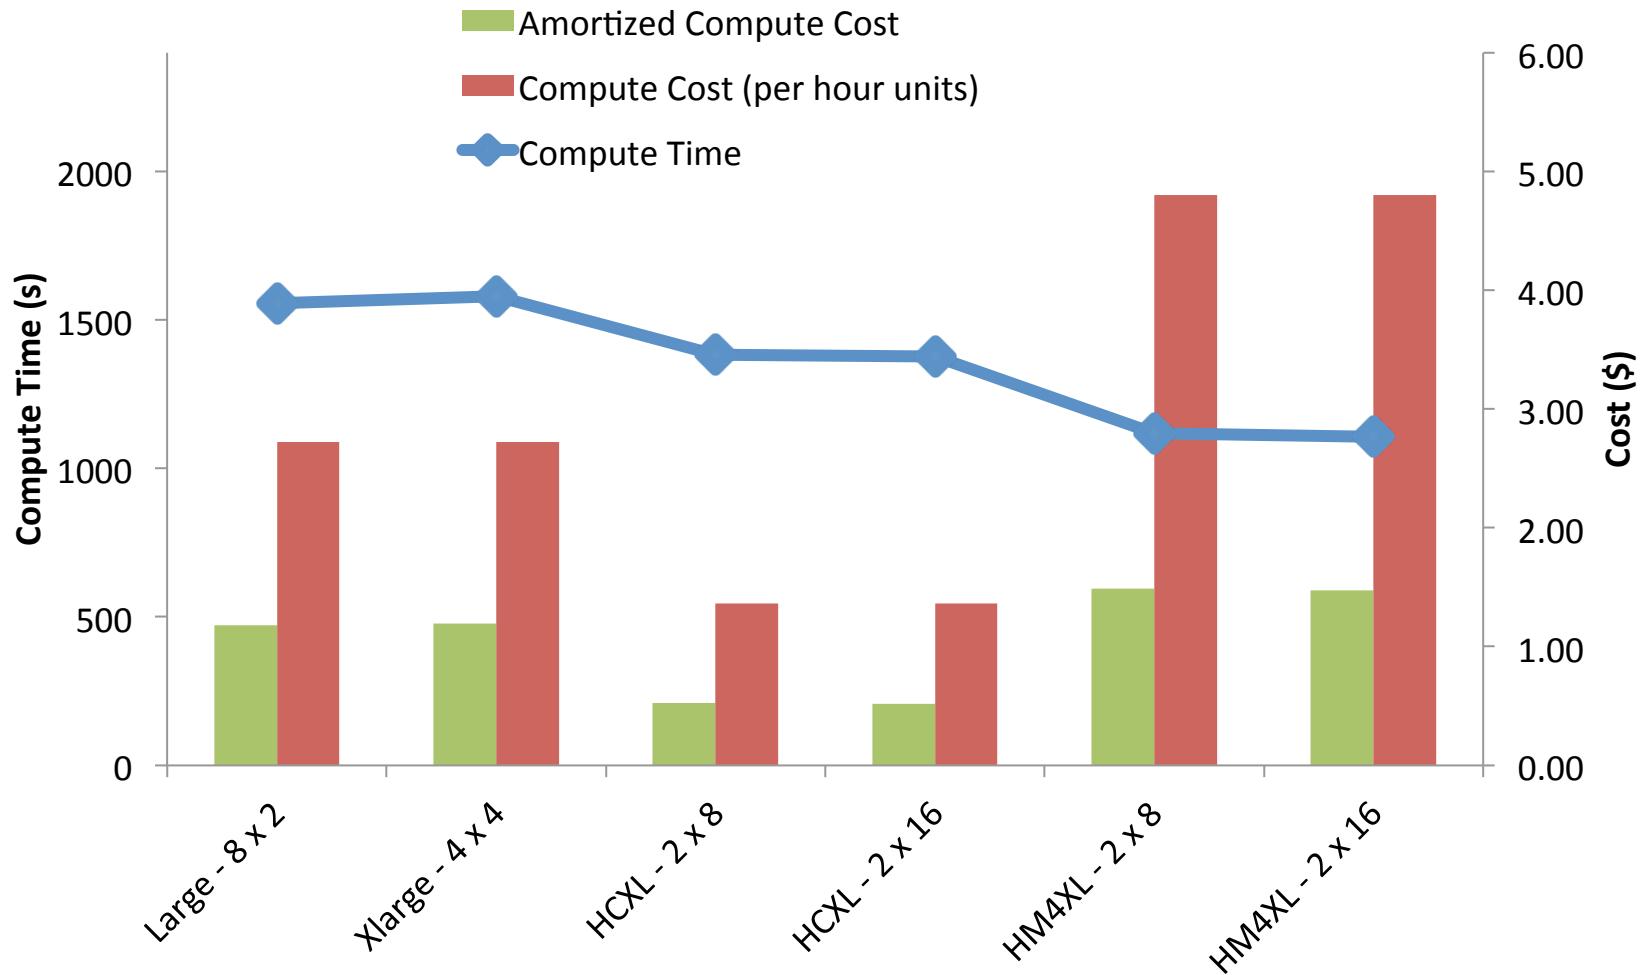
- Amazon Machine Image (AMI) is a special type of pre-configured operating system and virtual application software which is used to create a VM within EC2
 - Use either Pre-configured, templated images or create AMI to store customized images. Can share AMI (via AMI ID)
 - It serves as the basic unit of deployment for services delivered using EC2. Lease Linux as well as Windows AMI
 - See <http://aws.amazon.com/amazon-linux-ami/>
- VM = Bind an AMI to an Instance
 - Multiple Instance Types (see next slide)
 - Dynamically scale up/down
 - ‘root’ access to VM’s

Elastic Compute Cloud (EC2) Service

- EC2 Instances Types
 - <http://aws.amazon.com/ec2/instance-types/>
 - Standard Instance
 - Small, Large and Extra-Large
 - Micro Instance
 - High-Memory Instances
 - XL, Double XL, Quadruple XL
 - High-CPU Instances
 - High-CPU Medium, High-CPU XL
 - Cluster Compute Instance
 - Cluster Compute Quadruple
 - Cluster GPU Instance
 - ...

Sequence Assembly Performance with different EC2 Instance Types

(Adatped From Geoffrey Fox)



Azure

- Description: Microsoft’s “Platform as a Service” (PaaS) offering
 - Platform that is “Available” and “Scalable”
 - Cloud Based around virtualization
- Explicit Cost to Use
 - No cost to transfer data, only to use/store
- “Democratization of Infrastructure”
- Rich Data Abstractions
 - Large user data items: blobs
 - Service state: tables
 - Service workflow: queues
 - Simple and Familiar Programming Interfaces
 - REST: HTTP and HTTPS

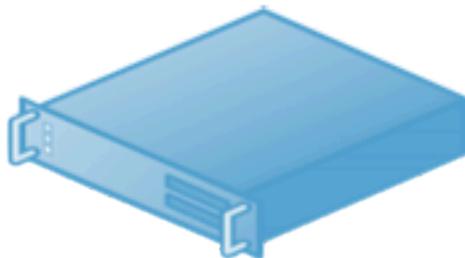
Each VM Has...

Processor

- CPU: Intel Xeon E5420
- Memory: 4 GB
- Network: 1000 Mbit/s
- Local Storage: 500 GB

Memory

- CPU: 8 cores
- Memory: 142 GB
- Local Storage: 27 TB

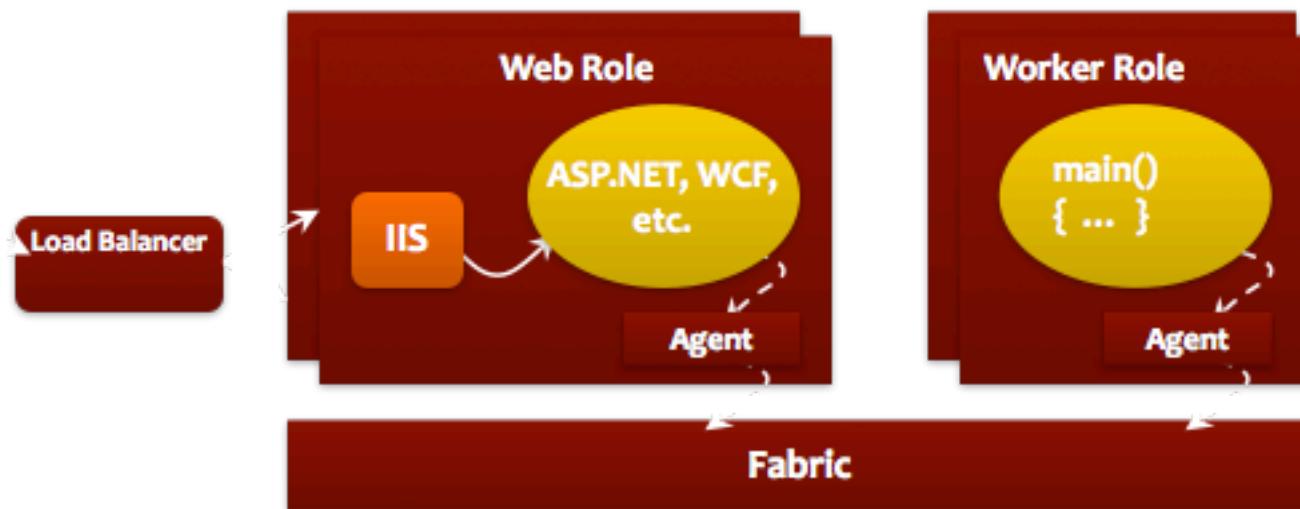




Windows Azure Compute Service

A closer look

HTTP



WAS

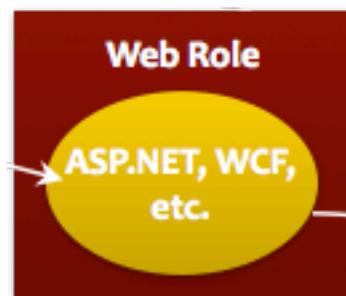


Suggested Application Model

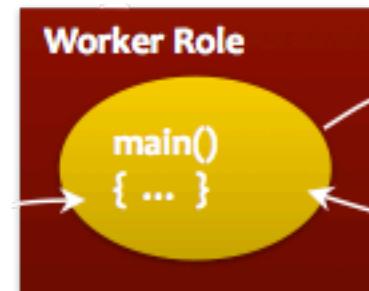
Using queues for reliable messaging

(two service model based on queue)

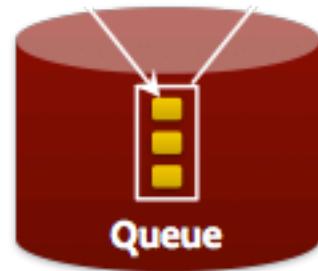
a) Receive work



b) Put work in queue

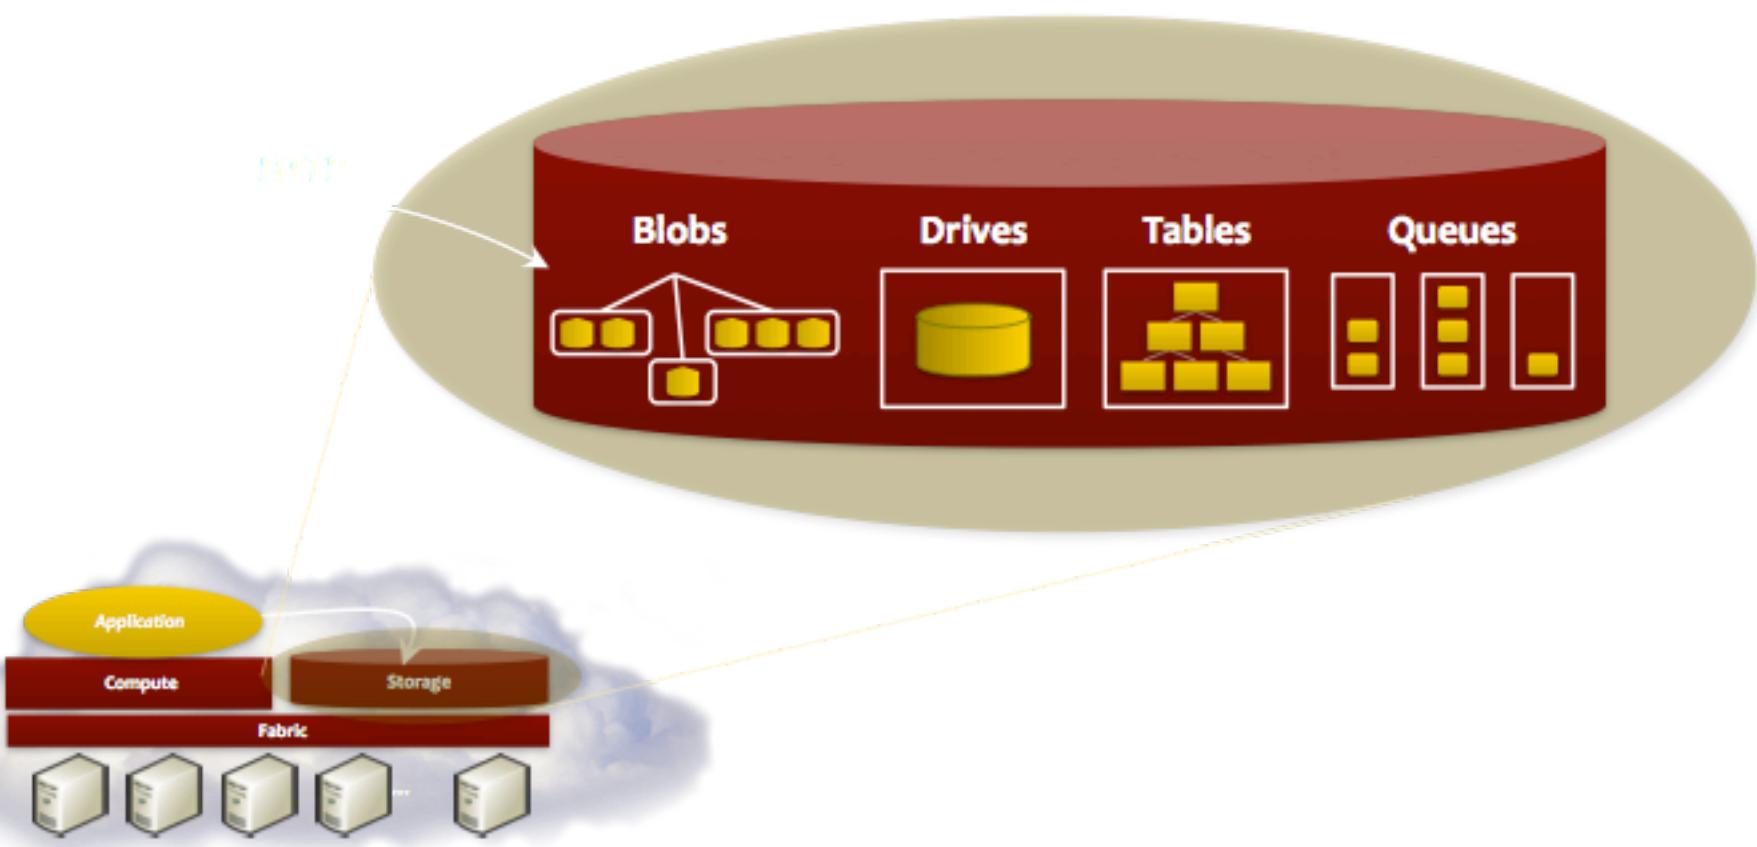


c) Get work from queue

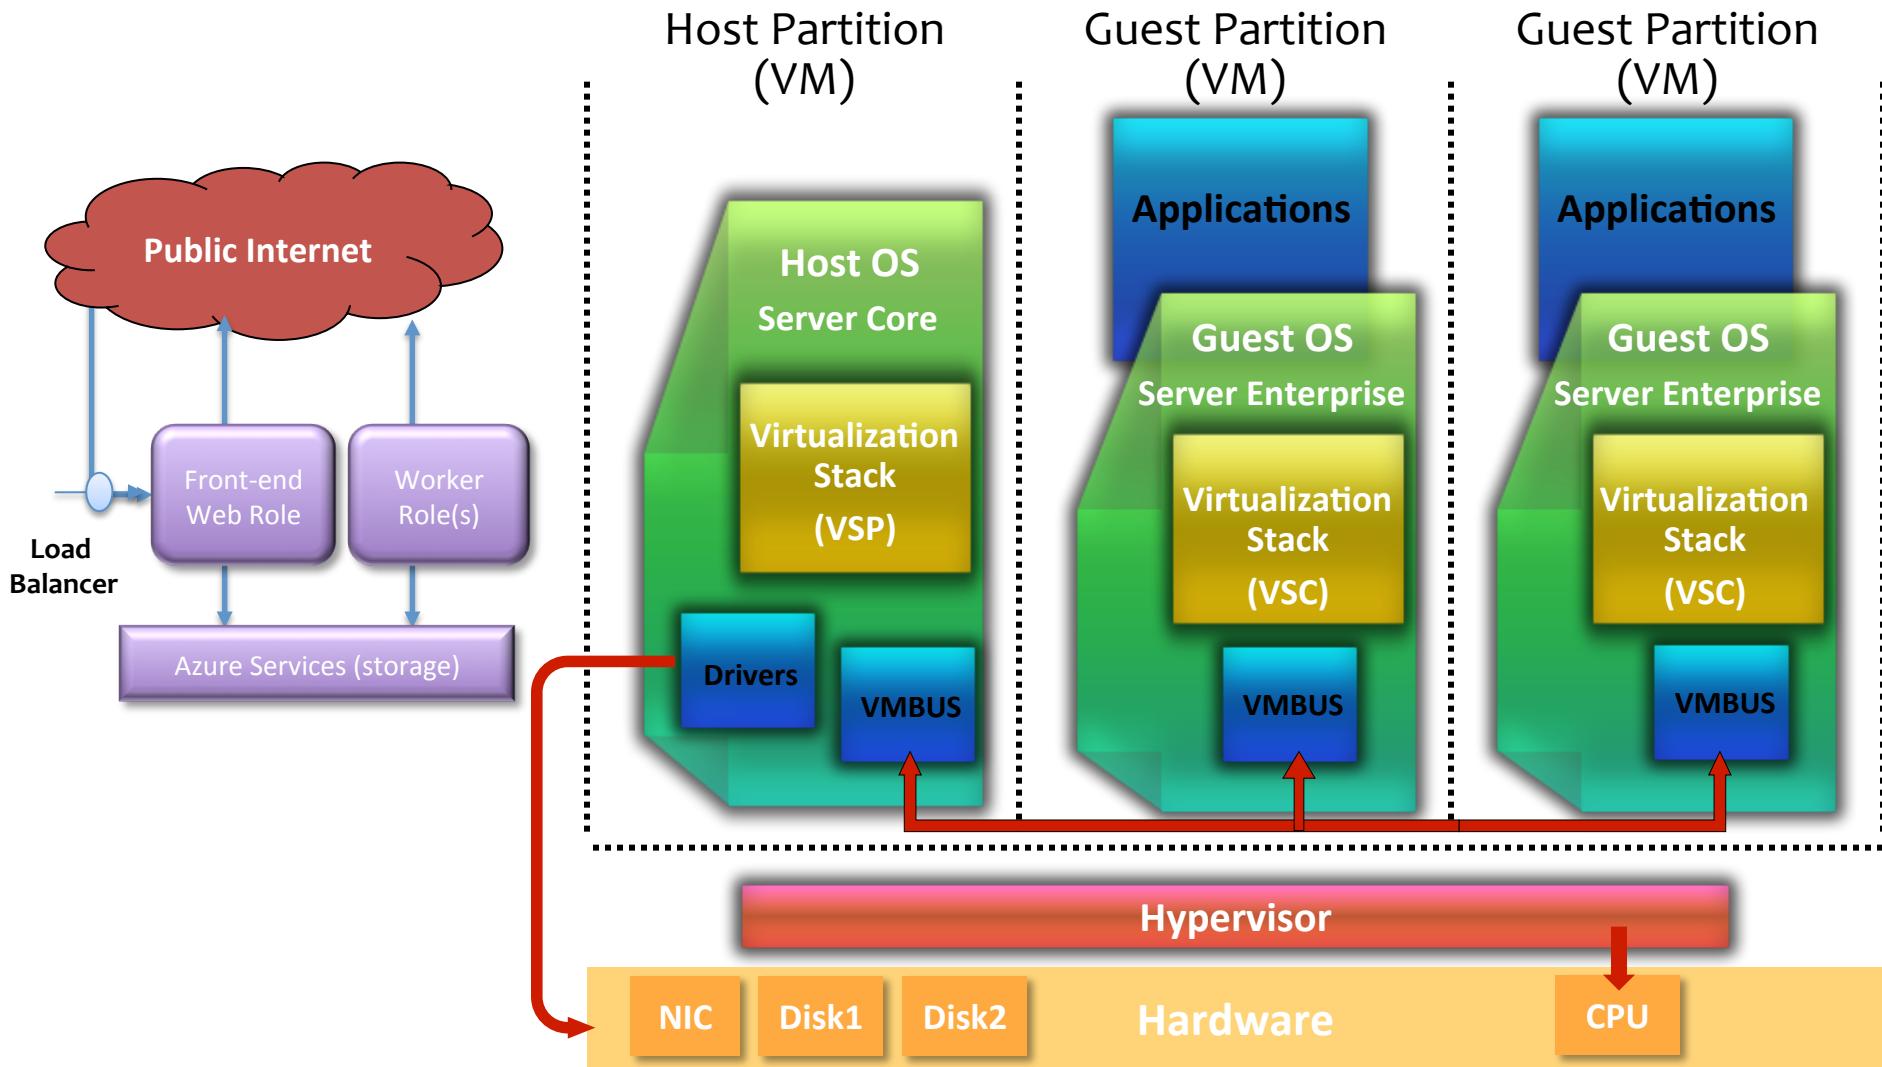


Azure Storage Service

A closer look



Azure Virtualization Architecture



Module E: Distributed Scientific Computing

To Distribute or not to Distribute?

Distributed Applications Summary

	Why Distributed?	How Distributed?	Challenges & Issues	How different from ?
Montage	Processing > local limits	Workflow enactor	Coordination	[1, 2]
NeKTAR	Processing > local limits (memory)	MPig	Advanced/Co-reservation	[1?, 4]
Ensemble-based/RE	Many compute-intensive task	SAGA, “Advert”	Coordination	[2,3]
ClimatePrediction.net	Many small tasks	BOINC, Trickles	Failures, variable # workers	[1, 4]
SCOOP	Peak req., naturally, Economic	Customized workflows	Not robust, adv. reservations	[1, 3, 4]

“Observations” on Distributed Applications

- Is large (and rich), but the number of effective and extensible DA small
 - More than just submitting jobs here and there!
- Developing DA is a hard undertaking
 - Intrinsic and Extrinsic Factors
 - Unique role of the Execution Environment (Infrastructure)
- Embrace “distributedness”
 - Understanding distributedness, heterogeneity & dynamic execution is fundamental (e.g., Exascale logically distributed prog. Models)
 - Data-centric application will be the drivers!
- Role for Pattern-oriented and Abstractions-based Development

Assertion #1: The space of possible DA is large, but number of effective DA small

- Distributed Application: That need multiple resources, or can benefit from the use of multiple resources;
 - .. can benefit from increased peak performance, throughput, reduced time-to-solution
 - More than just HPC or HTC Applications
 - e.g., DDDAS scenarios
- Ability to develop simple or effective distributed applications is limited
 - Applications that utilize multiple resources sequentially, concurrently or asynchronously is low
- Developing DA > just submitting jobs to remote sites!
 - What the pieces of distribution are? How these pieces interact? Flow of information? What is needed to actually deploy and execute the application?

Assertion #2: Developing DA is a hard undertaking

- Intrinsic reasons why developing DA is fundamentally hard:
 - Control & Coordination over Multiple & Distributed sites
 - Effective coordination in order for whole > sum of the parts
 - Complex design points; wide-range of models of DA
 - Many reasons for using DA, more than (just) peak performance
- Extrinsic:
 - Execution environments will be dynamic, heterogeneous and varying degrees-of-control
 - Fundamental different variation in role of Execution Environment- distinguishing feature of DA from “regular environment” HPC
 - Application types strongly coupled to the infrastructure capabilities, abstractions/tools, & policy:
 - Often development tools assume “specific” deployment and execution environments, or don’t where needed!
 - Policies and tools, e.g production DCI has been missing for DDDAS

Assertion #2: Developing DA is a hard undertaking

- Large number programming systems, tools and environments
 - Lack of extensible functionality, interfaces & abstractions
 - Interoperability and extensibility become difficult
 - *Art of tool building needs to be more of science!*
- Applications have been brittle and not extensible:
 - Tied to specific tools and/or programming system
 - *Large number of Incomplete Solutions!*
- Unique Role for abstractions for DA and CI
 - Application formulation, development and execution must be less dependent on infrastructure & provisioning details
 - Abstractions for Development, Deployment & Execution
 - A Pattern-Oriented, Abstractions-Based Approach
 - “Abstractions allows innovation at more interesting layers”

Assertion #3: Embrace Distribution

- “History of computing like pendulum, swings from centralized to distributed”
 - Indications this time there is a fundamental paradigm shift due to DATA
 - Too much to move around; learn how to do analytics/compute *in situ*
- Decoupling and delocalization of the producers-consumers of computation
 - Localized special services; people and collaborations are distributed
- (Ironically) Most applications have been developed to hide from heterogeneity and dynamism; not embrace them
 - Programming models that provide dynamic execution (opposed to static), address heterogeneity etc
 - Logically vs Physically Distributed: NG programming models will need to support dynamic execution, heterogeneity at a logically-distributed level

Assertion #3: Embrace Distributedness

Corollary: Clouds are not Panacea

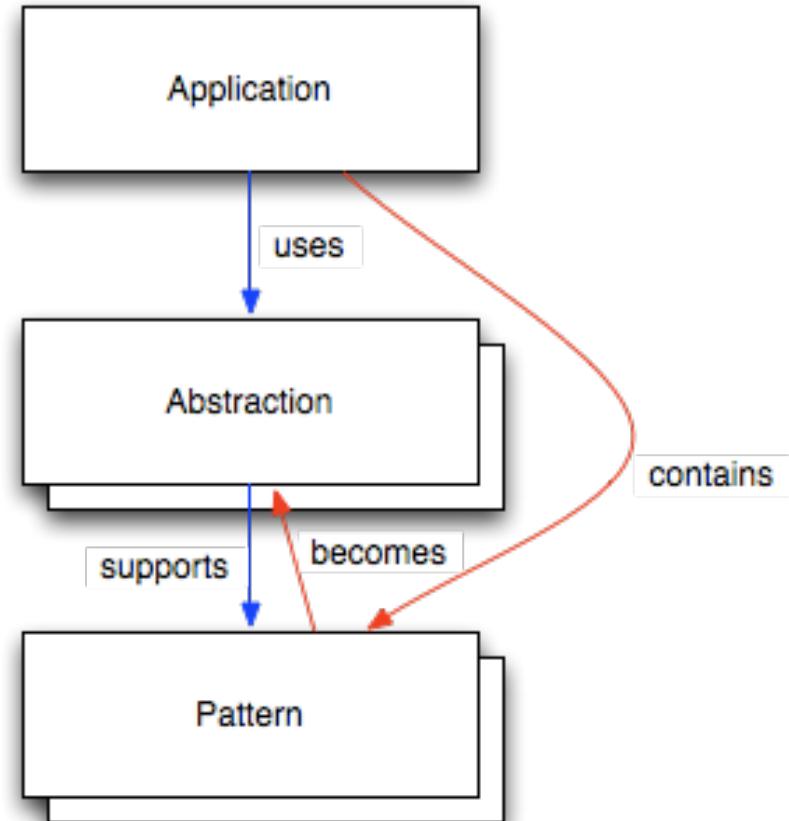
- Clouds: Novel or more of the same?
 - Better control over software environment via virtualization
 - Illusion of unlimited and immediate available resource can lead to better capacity planning and scheduling
 - Partly due to underlying economic model and SLAs
- Clouds do not remove many/all of the challenges inherent in DA
 - Clouds are about provisioning, grids are about federation
 - Fundamental challenges in distribution remain
 - Makes some thing worse as impose a model of strong localization
 - *"The reason why we are so well prepared to handle the multi-core era, is because we took the trouble to understand and learn parallel programming"* – Ken Kennedy
- Clouds part of a larger distributed CI
 - Certain tasks better suited for Grids, others on Clouds

Assertion #4: Role for a Pattern-Oriented and Abstraction-Based Development Cycle

■ Relation between Application, Abstractions and Patterns:

- Application: Need or can use >1 R
- Patterns: Formalizations of commonly occurring modes of computation, composition, and/or resource usage
 - Devel, Deploy & Exec Phase
- Abstractions: Process, mechanism or infrastructure to support a commonly occurring usage

Coordination	Deployment
Master-Worker (TF, BoT)	Replication
All-Pairs	Co-allocation
Data Processing Pipeline	Consensus
MapReduce	Brokeraging
AtHome	
Pub-Sub	
Stream	



Assertion #4: Role for a Pattern-Oriented and Abstraction-Based Development Cycle

- Analysis of Distributed Applications leads to three types of patterns
 - Patterns that appear in the Parallel Programming
 - Patterns driven by distributed concerns (eg @HOME, consensus)
 - Patterns addressing distributed environment concerns exclusively (eg co-allocation)
- There exists tools that support patterns, i.e., provide abstractions

Pattern	Tools That Support the Pattern
Master/Worker-TaskFarm	Aneka, Nimrod, Condor, Symphony, SGE, HPCS
Master/Worker-BagofTasks	Comet-G, TaskSpace, Condor, TSpaces
All-Pairs	All-Pairs
Data Processing Pipeline	Pegasus/DAGMan
MapReduce	Hadoop, Twister, Pydoop
AtHome	BOINC
Pub-Sub	Flaps, Meteor, Narada, Gryphon, Sienna
Stream	DART, DataTurbine
Replication	Giggle, Storm, BitDew, BOINC
Co-allocation	HARC, GUR
Consensus	BOINC, Chubby, ZooKeeper
Brokers	GridBus, Condor matchmaker

Application Example	Coordination	Deployment
Montage	TaskFarm, Data Processing Pipeline	-
NEKTAR	-	Co-allocation
Coupled Fusion Simulation	Stream	Co-allocation
Async RE	Pub/Sub	Replication
Climate-Prediction (generation)	Master/Worker, AtHome	Consensus
Climate-Prediction (analysis)	MapReduce	-
SCOOP	Master/Worker, Data Processing Pipeline	-

IDEAS: DA Development Objectives

- **Interoperable:** Ability to work across multiple resources concurrently
 - Includes jobs submission, coordination mechanism,
- **Dynamic:** Beyond legacy static execution & resource allocation models
 - Decisions at both deployment and run-time
 - Dynamical execution is almost fundamental at scale
- **Extensible:** Support new functionality & infrastructure without wholesale refactoring, i.e., lower coupling to tools & infrastructure
- **Adaptive/Autonomic:** Flexible response to fluctuations in dynamic resources, availability of dynamic data
- **Scalable:** Along many dimensions and design points

Challenge: To develop DA effectively and efficiently with IDEAS as first class objectives with simplicity an over-aching concern

Module E: Project Redux

- Gain sufficient proficiency with SAGA to write a M-W (from scratch) application that uses > 1 XSEDE resource?
- Use Clouds:
 - Can use SAGA to submit jobs to FG-based Clouds?
 - Compare Application X on XSEDE on Clouds?
- ...
- *Teamwork is acceptable provided: (i) effort is acknowledged, (ii) clear intellectual contribution from each*

References

- Python-based Master-Worker:
 - <http://pymw.sourceforge.net/>
- Google MapReduce
 - <http://code.google.com/edu/parallel/mapreduce-tutorial.html>
- <http://groups.google.com/group/vscse-big-data-for-science-2010/web/course-presentations>
- <http://futuregrid.org/tutorials>

M-W: Issues to consider

- https://svn.cct.lsu.edu/repos/saga-projects/applications/master_worker
- Aim: Understand trade-off issues along three dimensions:
 - (i) work decomposition (ii) distribution and (iii) coordination
- Homework:
 1. Everything local: For 1 Master and same workload vary: $N_w = 2, 4$ and 8 Plot times to completion.
 2. With the advert service running remotely, repeat the above. Compare performance with (1)
 3. With the advert service running distributed: Distribute (equally?) the workers across a couple of FutureGrid machines. Compare with (i) and (2)
 4. Extend with user defined “Worker”.. use simple worker function.